

Olmir Stocker Neto

AI ENGINEER · LLM SYSTEMS · MULTI-AGENT ARCHITECTURES

📍 São Paulo, SP, Brazil

📞 +55 11 94518-6180

✉ owmyrstocker@gmail.com

🌐 linkedin.com/in/olmir-stocker-neto

🐙 github.com/owmyr

🌐 owmyr.github.io

PROFESSIONAL SUMMARY

Software Engineer specializing in applied Generative AI, autonomous multi-agent architectures, and production LLM integrations (Python, TypeScript). Proven track record of shipping resilient AI systems—including enterprise knowledge assistants at Accenture, fault-tolerant agent swarms, and hybrid RAG pipelines with vector reranking. Experienced in turning non-deterministic foundation models into secure, reliable, and observable software.

TECHNICAL SKILLS

Programming Languages:

Python, TypeScript, JavaScript (ESNext), SQL

AI & Machine Learning:

LLM Orchestration, Multi-Agent Systems, RAG, Prompt Engineering, FAISS (Vector DBs), Cross-Encoder Reranking, PyTorch, Scikit-learn, LangChain, Hugging Face

Backend & Systems:

FastAPI, Node.js, Zod, Pydantic, PostgreSQL, SQLite (WAL), Firebase/Firestore, REST APIs

DevOps, Testing & Tooling:

AWS (SageMaker), Docker, GitHub Actions (CI/CD), Vitest, pytest, Git, Agile/Scrum

PROFESSIONAL EXPERIENCE

AI Engineer (Freelancer Contract)

Oct 2025 – Present

Revelo · Freelance / Project-Based · Remote

- Applied advanced prompt engineering to diagnose and document failure modes in generative AI models for code generation tasks.
- Acted as a Human-in-the-Loop (HITL) reviewer, evaluating **100+ code snippets** to guarantee production quality and resolving critical model degradation issues.
- Authored **50+ unit tests** with `pytest` to validate deterministic correctness and execution safety of AI-generated solutions.
- Managed and replicated complex multi-dependency runtime environments from GitHub commits using AWS SageMaker and Docker.
- Delivered **200+ approved tasks** across multiple AI engineering workflows, completing **50+ peer reviews** to improve model alignment.

AI Agent Evaluation Engineer

May 2026 – Sep 2026

Turing · Contract · Remote

- Engineered and evaluated multi-step **Chain-of-Thought (CoT)** reasoning trajectories for frontier LLMs, formalizing step-by-step mathematical reasoning and optimal algorithmic trade-offs.
- Authored complex algorithmic reasoning benchmarks with LaTeX mathematical formulations, enforcing deterministic output constraints and eliminating model hallucinations.
- Containerized automated evaluation harnesses within **Docker sandboxes**, executing **30+ edge-case test suites** per task to stress-test model reasoning and verify time/memory complexity boundaries.

Software Engineer

Jan 2026 – May 2026

Accenture · Hybrid · Enterprise Client

- Deployed an enterprise AI Knowledge Assistant utilizing **Microsoft Copilot** and **RAG architecture**, designing specialized domain sub-agents that reduced developer onboarding time by 40% and saved 100+ senior engineering hours across enterprise configuration workflows.
- Engineered high-volume ETL pipelines and complex SQL extraction queries integrated directly with SAP IBP to process nationwide datasets and support operational planning.
- Contributed to the modernization of legacy supply chain architecture into a scalable, high-throughput distributed ecosystem.
- Collaborated in an Agile Scrum framework to map API contracts, author technical specifications, and ensure cross-team alignment.

FEATURED AI ENGINEERING PROJECTS

Multi-Agent Swarm AI Orchestration Layer

github.com/owmyr/care-agent-swarm

Technologies: TypeScript, Node.js, Claude 3.5 Sonnet, Zod Contracts, Multi-Agent Swarms, Circuit Breaker, HIPAA Safe Harbor, Vitest (42 Tests)

- Architected a multi-agent orchestration layer for residential care workflows, automating resident intake, compliance checks, and incident triage.
- Implemented concurrent sub-agents with fault-isolated execution, allowing partial intake completion even if an individual agent experiences transient degradation.
- Engineered a resilient LLM harness featuring a circuit breaker state machine, exponential backoff, and automatic handling of provider rate limits (HTTP 429).
- Enforced strict runtime data contracts using Zod schemas to validate and re-prompt malformed model outputs, guaranteeing deterministic JSON payloads.
- Integrated HIPAA-compliant PHI redaction across application logs and built an automated clinical escalation safeguard for unresolved incident loops.
- Authored 40+ automated unit and integration tests with Vitest using mock client dependency injection for deterministic, zero-cost test execution.

TrendScout — Autonomous Trend Intelligence & Multi-Agent Pipeline

github.com/owmyr/TrendScout · trendscout-shopee.vercel.app

Technologies: Python 3.11+, Google Gemini Flash, Local Ollama, Multi-Agent Fleet, Playwright, Next.js 15, Air-Gapped Pipeline, pytest (222 Tests)

- Engineered an end-to-end intelligence platform that tracks e-commerce sales velocity, eliminates commodity noise, and semantically clusters apparel design trends.
- Built a dual-tier LLM engine utilizing a Google Gemini Flash pool for high-throughput clustering with seamless offline fallback to a local Ollama model.
- Designed an adaptive discovery algorithm combining historical baseline keywords with velocity-driven search terms to catch emerging micro-trends early.
- Implemented deterministic regex pre-filtering to remove plain/basic garments prior to LLM inference, reducing token consumption while preserving audit baselines.
- Built an air-gapped data export pipeline feeding a Next.js 15 subscriber portal, backed by a comprehensive suite of 220+ automated pytest tests.

Hybrid RAG Chatbot with Cross-Encoder Re-Ranking

github.com/owmyr/D-D-RAG-Chatbot

Technologies: Python, LangChain (LCEL), FAISS Vector Index, TinyBERT Cross-Encoder, Gemini API, Header-Aware Chunking

- Architected a modular Hybrid RAG pipeline using LangChain and Gemini to query complex documentation with sub-second retrieval latency.
- Built an ingestion pipeline with markdown header-aware chunking to preserve document hierarchy, breadcrumbs, and cross-references.
- Implemented a two-stage retrieval pipeline combining FAISS vector search with a local Cross-Encoder (TinyBERT) reranker for high contextual precision.
- Engineered atomic index swapping using staging directories to prevent vector store corruption during updates, cutting indexing time by 50%.

EDUCATION & CERTIFICATIONS

Postgraduate Specializations · Educaminas

2025 – 2027 (Expected)

- **Software Architecture:** Microservices, SOA, Distributed Systems, Cloud Architecture, DevOps
- **Data Science & Machine Learning:** Big Data (Hadoop, Spark), Deep Learning, Predictive Analytics
- **Applied Statistics:** Demand Forecasting, ANOVA, Experimental Design, Financial Mathematics

Bachelor in Systems Analysis and Development

2021 – 2024

UNIP · São Paulo, Brazil

AI Professional Training · Hashtag Treinamentos

Practical, project-based programming covering scalable Python architectures, data modeling, and machine learning pipelines.

LANGUAGES

English: C2 Level (Fluent · Full Professional Proficiency) Portuguese: Native